

2nd Hand Car Price Prediction

In this project we analyze the 'Used Cars' dataset from kaggle (<https://www.kaggle.com/datasets/austinreese/craigslist-carstrucks-data?select=vehicles.csv>).

The goal is to build a model that could estimate the price of second hand cars based on relevant features. We explore the data through cleanig and preprocessing, handling outliers, feature engineering and eventually testing different regression models to get the most accurate predictor.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib
from matplotlib import pyplot as plt
import seaborn as sns
%matplotlib inline
matplotlib.rcParams['figure.figsize'] = (20,10)
```

```
In [2]: df1 = pd.read_csv("/Users/marwa/Desktop/2ndHCP/model/vehicles.csv")
```

Data Cleaning

```
In [4]: df1.head()
```

```
Out [4]:
```

	id	url	region	
0	7222695916	https://prescott.craigslist.org/cto/d/prescott...	prescott	https://prescott.craigslist.org/cto/d/prescott...
1	7218891961	https://fayar.craigslist.org/ctd/d/bentonville...	fayetteville	https://fayar.craigslist.org/ctd/d/bentonville...
2	7221797935	https://keys.craigslist.org/cto/d/summerland-k...	florida keys	https://keys.craigslist.org/cto/d/summerland-k...
3	7222270760	https://worcester.craigslist.org/cto/d/west-br...	worcester / central MA	https://worcester.craigslist.org/cto/d/west-br...
4	7210384030	https://greensboro.craigslist.org/cto/d/trinit...	greensboro	https://greensboro.craigslist.org/cto/d/trinit...

5 rows x 26 columns

```
In [5]: df1.shape
```

```
Out [5]: (426880, 26)
```

```
In [9]: df1.describe()
```

Out [9]:

	id	price	year	odometer	county	
count	4.268800e+05	4.268800e+05	425675.000000	4.224800e+05	0.0	42033
mean	7.311487e+09	7.519903e+04	2011.235191	9.804333e+04	NaN	3
std	4.473170e+06	1.218228e+07	9.452120	2.138815e+05	NaN	
min	7.207408e+09	0.000000e+00	1900.000000	0.000000e+00	NaN	-8
25%	7.308143e+09	5.900000e+03	2008.000000	3.770400e+04	NaN	3
50%	7.312621e+09	1.395000e+04	2013.000000	8.554800e+04	NaN	3
75%	7.315254e+09	2.648575e+04	2017.000000	1.335425e+05	NaN	4
max	7.317101e+09	3.736929e+09	2022.000000	1.000000e+07	NaN	8

Drop irrelevant features:

Region can be removed since state will suffice

Note that county has no values

Description and type can be removed since we are relying on manufacturer

```
In [11]: df2 = df1.drop(['id', 'url', 'region_url', 'image_url', 'description', 'la
                'long', 'region', 'VIN', 'title_status', 'type', 'cylinders
```

Handling missing values

```
In [13]: #check columns with nan>50%
```

```
df2.isnull().mean()*100
```

```
Out[13]: price          0.000000
year          0.282281
manufacturer    4.133714
model          1.236179
condition      40.785232
fuel           0.705819
odometer       1.030735
transmission    0.598763
drive          30.586347
size           71.767476
paint_color     30.501078
state          0.000000
posting_date    0.015930
dtype: float64
```

```
In [14]: #since size feature as 71% NaN we will remove this feature
#and since our dataset is large it is efficient enough to remove all rows
df2 = df2.drop('size', axis='columns')
df3 = df2.dropna()
df3.shape
```

```
Out[14]: (152629, 12)
```

Feature Engineering

Checking categories of categorical variables

```
In [343]: print("\nmanufacturer:\n" , len(df3['manufacturer'].unique()),
            "\nmodel:\n", len(df3['model'].unique()),
            "\ncondition:\n", len(df3['condition'].unique()),
            "\nfuel:\n", len(df3['fuel'].unique()),
            "\n transmission:\n", len(df3['transmission'].unique()),
            "\n drive: \n", len(df3['drive'].unique()),
            "\n paintcolour\n", len(df3['paint_color'].unique()),
            "\n state: \n", len(df3['state'].unique()) )
```

```
manufacturer:
  41
model:
 12963
condition:
   6
fuel:
   5
transmission:
   3
drive:
   3
paintcolour
  12
state:
  51
```

```
In [18]: #since model has a lot of different categories it can make interpretability
#we will remove it since we assume that manufacturer provides enough info
```

```
df4=df3.drop(['model'], axis='columns')
```

```
In [19]: #we will remove 'harley-davidson' from manufacturer as it is a motorcycle
```

```
df4= df4[~(df4.manufacturer== 'harley-davidson')]
```

Checking count and reduce if necessary to prevent the dimensionality curse

that arises when performing one-hot-encoding in high dimensional categorical variables

what we will do is check how many samples are present per category for all manufacturer,

paint color and state

```
In [20]: manufacturer_count= df4.groupby('manufacturer')['manufacturer'].agg('count')
paint_count=df4.groupby('paint_color')['paint_color'].agg('count').sort_values(ascending=True)
state_count= df4.groupby('state')['state'].agg('count').sort_values(ascending=True)

print(manufacturer_count,
```

```
'\n\n ', paint_count ,  
'\n\n ', state_count)
```

```
manufacturer
ford          26172
chevrolet     21935
toyota        12902
honda         8599
nissan        7705
jeep          7345
gmc           5786
bmw           5223
dodge         5115
ram           4920
volkswagen    4293
hyundai       4030
mercedes-benz 3853
subaru        3340
kia           3206
lexus         3020
mazda         2382
chrysler      2368
cadillac      2365
buick         2163
acura         1977
lincoln       1828
infiniti      1811
audi          1755
mitsubishi    1450
volvo         1122
pontiac       1066
mini          1035
rover         681
jaguar        585
mercury       563
saturn        526
tesla         452
porsche       434
fiat          316
alfa-romeo    199
ferrari       25
datsun        24
land rover    9
aston-martin  8
Name: manufacturer, dtype: int64
```

```
paint_color
white      37736
black      30835
silver     21870
blue       16865
red        16835
grey       14621
green      4473
brown      4035
custom     2759
yellow     1177
orange     980
purple     402
Name: paint_color, dtype: int64
```

```
state
ca      17824
```

fl	10208
ny	8711
tx	7670
oh	7010
mi	6160
pa	5702
nc	5397
wi	5048
ma	3897
tn	3871
va	3665
or	3656
il	3641
co	3558
nj	3528
ia	3455
az	3285
mn	3274
in	2768
ok	2706
ga	2535
ks	2398
sc	2358
id	2241
ct	2062
ky	2056
wa	1867
nm	1698
mo	1677
al	1654
md	1530
vt	1512
ar	1497
mt	1423
nh	1171
ri	1138
me	1121
dc	1033
ak	1022
nv	965
la	912
hi	664
sd	505
de	444
wv	429
ms	402
ne	387
ut	354
wy	269
nd	230

Name: state, dtype: int64

```
In [21]: manufacturer_less_than_100 = manufacturer_count[manufacturer_count<100]
len(df4['manufacturer'].unique())
```

Out[21]: 40

```
In [22]: # we will group the manufacturers with counts <100 together as 'other'
```

```
df4.manufacturer= df4.manufacturer.apply(lambda x: 'other' if x in manufa
len(df4['manufacturer'].unique())
```

Out [22]: 37

Extracting age of each vehicle

```
In [24]: # want the feature post_date to be a year only

#first we convert it to a date time for easier manipulation
#using to_datetime from pandas and setting utc=True for correct handling
df4['posting_date'] = pd.to_datetime(df4['posting_date'], utc=True)

df4['posting_date'].dtypes
```

Out [24]: datetime64[ns, UTC]

```
In [25]: #then we extract the year out of it using dt.year and create a new column
df4['posting_year'] = df4['posting_date'].dt.year

#remove postiing date column
df4= df4.drop('posting_date', axis= 'columns')
df4.year.dtypes
```

Out [25]: dtype('float64')

```
In [26]: #turn year into int32 to allow subtraction
df4['year']= df4['year'].astype('int32')
```

```
In [27]: #now we want to create a column of age of the car since purchase

df5 = df4.copy()
df5['vehicle_age'] = df5['posting_year']-df5['year']
```

```
In [28]: #clearly vehicle_age and year now are highly correlated and the posting_y
#2021 for all its entries which does not provide much information
#we will keep vehicle age as it provides enough info about the 2

df5=df5.drop(['year', 'posting_year'], axis= 'columns')
```



Outlier Removal

target variable price

```
In [32]: #Visualising target variable price

plt.figure(figsize=(12, 6))
sns.histplot(df5['price'], bins=50, kde=True)
plt.title('Histogram of Price Distribution')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```

```
plt.figure(figsize=(12, 6))
sns.boxplot(x=df5['price'])
plt.title('Box Plot of Price Distribution')
plt.show()
```



In [33]: *# check price range*

```
df5.price.describe(percentiles = [0.25,0.50,0.75,0.85,0.90,1])
```



```
Out [33]: count    1.525880e+05
          mean     7.598327e+04
          std      1.383998e+07
          min      0.000000e+00
          25%      6.000000e+03
          50%      1.290000e+04
          75%      2.399500e+04
          85%      3.099000e+04
          90%      3.499500e+04
          100%     3.736929e+09
          max      3.736929e+09
          Name: price, dtype: float64
```

mean price= 75,967 std dev= 13,838,120 which is extremely high relative to the mean.

This indicates that there's a wide spread in the price and the possible presence of outliers.

min=0 indicates that there're samples with no price!! this can't be marketable. max= 3.74 billion! is extremely high for a car price! Could be an outlier

First Remove cars with price=0 since they are not marketable and dont provide any insight to our analysis

```
In [36]: df6= df5[~(df5.price==0)]
          df6[df5['price']==0]
```

```
/var/folders/s3/24r6s08x3pg9xyqf7659_57h0000gn/T/ipykernel_88375/52173903
3.py:2: UserWarning: Boolean Series key will be reindexed to match DataFra
me index.
      df6[df5['price']==0]
```

```
Out [36]: price  manufacturer  condition  fuel  odometer  transmission  drive  paint_color
```

since prices vary in different states and the figures are highly

right skewed it is best to perform the IQ method on price per state,

so for every location we will get the bounds of the price in that

location then remove outliers for each location

```
In [40]: def remove_outliers(df):
          df_out= pd.DataFrame()
          #group by state
          for key, subdf in df.groupby('state'):

              Q1= subdf['price'].quantile(0.25)
              Q3= subdf['price'].quantile(0.75)
              IQR= Q3-Q1

              #define bounds
              lower= Q1 - 1.5 * IQR
              upper= Q3 + 1.5 * IQR
```

```
#select prices per state within the bounds
reduced_df= subdf[(subdf.price>= lower) & (subdf.price<= upper) ]
df_out= pd.concat([df_out, reduced_df], ignore_index=True)

return df_out

df7=remove_outliers(df6)
df7.shape
```


Out[40]: (144078, 10)

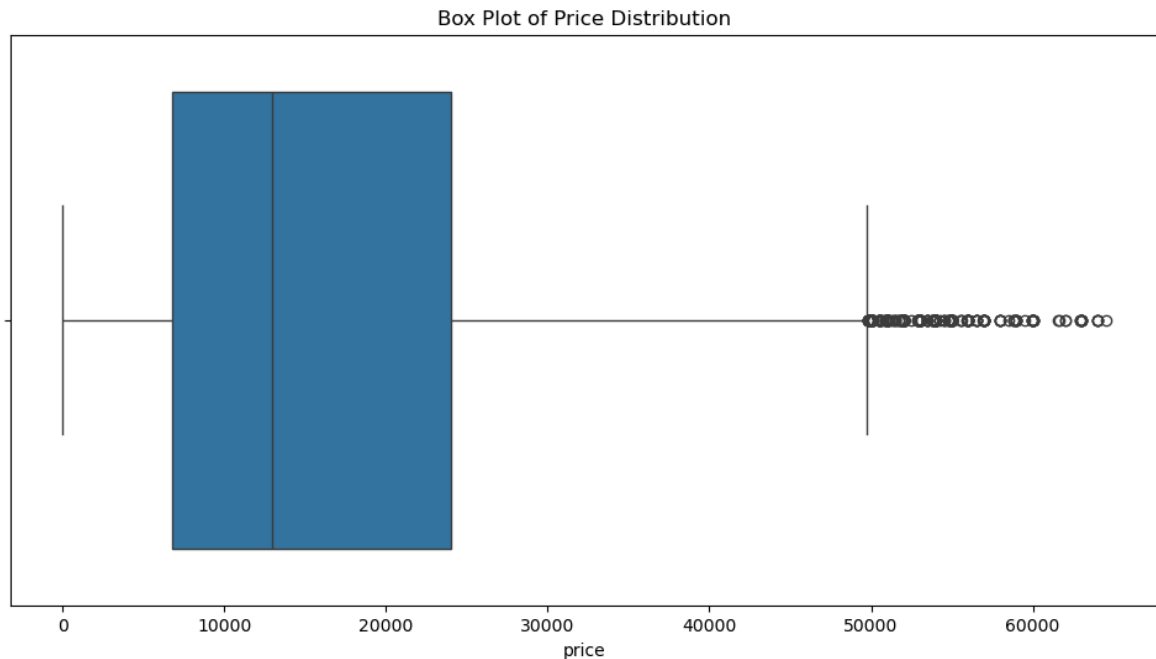
Price target variable

```
In [42]: # visualise Price again

plt.figure(figsize=(12, 6))
sns.histplot(df7['price'], bins=50, kde=True)
plt.title('Histogram of Price Distribution')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()

plt.figure(figsize=(12, 6))
sns.boxplot(x=df7['price'])
plt.title('Box Plot of Price Distribution')
plt.show()
```





```
In [43]: df7['price'].describe(percentiles = [0.25,0.50,0.75,0.85,0.90,1])
```

```
Out[43]: count      144078.000000
mean       16103.381821
std        11566.195793
min         1.000000
25%        6795.000000
50%       12995.000000
75%       23990.000000
85%       29990.000000
90%       33990.000000
100%      64500.000000
max       64500.000000
Name: price, dtype: float64
```

The range looks better now it is still right skewed but that is to be expected in car prices

odometer/ mileage feature

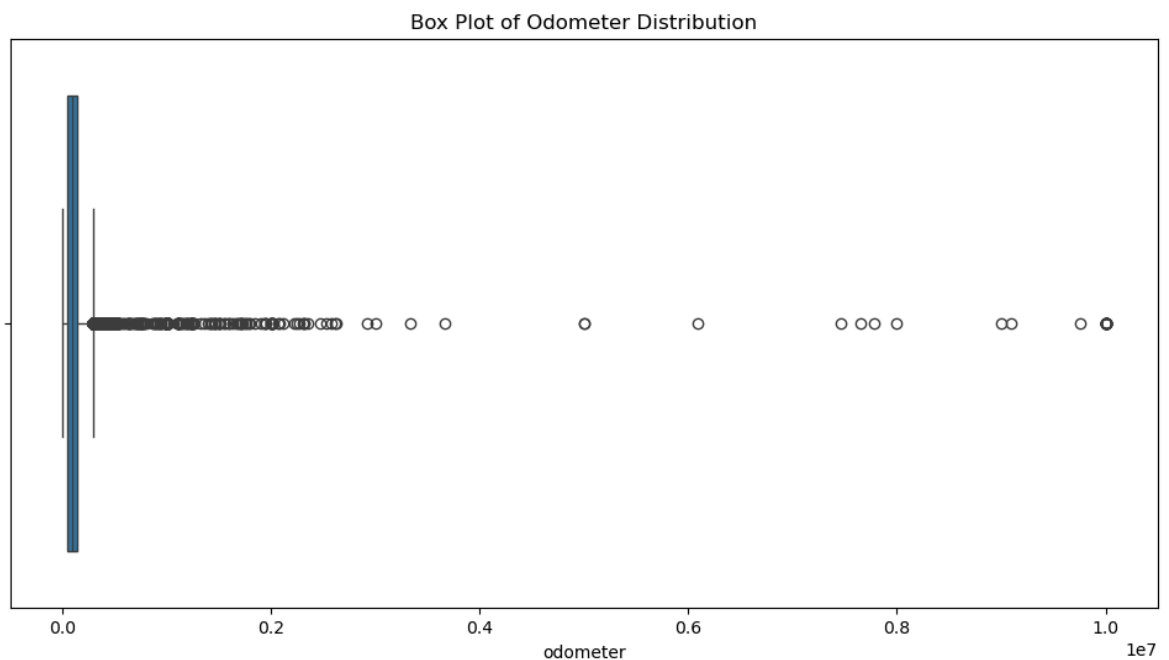
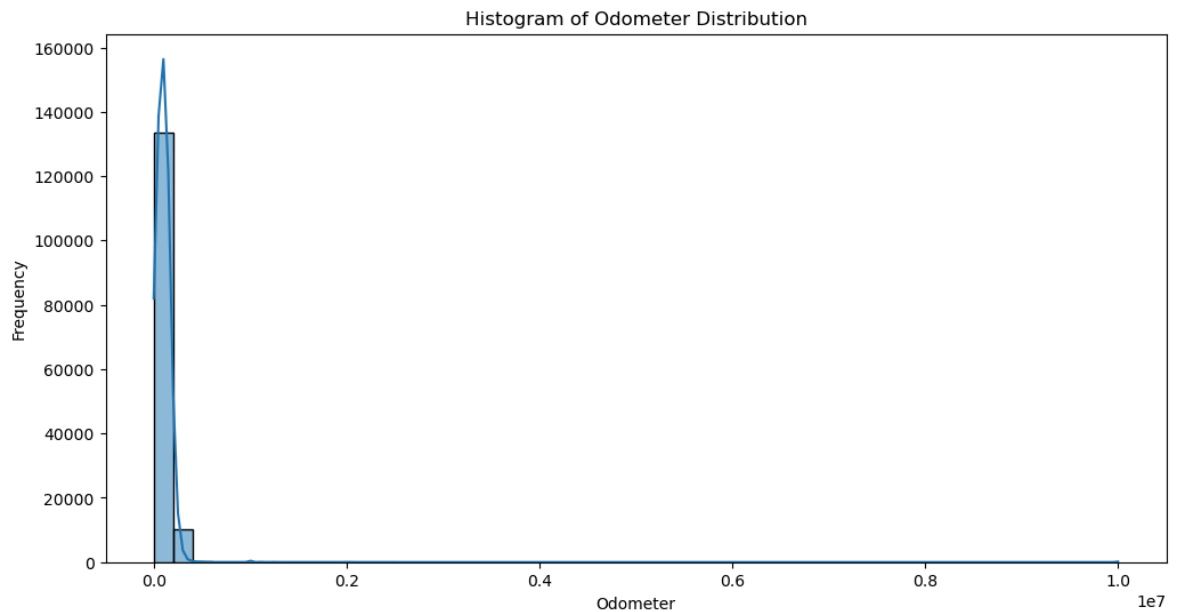
```
In [46]: df7['odometer'].describe(percentiles = [0.25,0.50,0.75,0.85,0.90,1])
```

```
Out[46]: count      1.440780e+05
mean       1.039443e+05
std        1.742401e+05
min        0.000000e+00
25%        4.397425e+04
50%        9.700000e+04
75%       1.435270e+05
85%       1.684450e+05
90%       1.860000e+05
100%      1.000000e+07
max       1.000000e+07
Name: odometer, dtype: float64
```

```
In [47]: plt.figure(figsize=(12, 6))
sns.histplot(df7['odometer'], bins=50, kde=True)
plt.title('Histogram of Odometer Distribution')
```

```
plt.xlabel('Odometer')
plt.ylabel('Frequency')
plt.show()

plt.figure(figsize=(12, 6))
sns.boxplot(x=df7['odometer'])
plt.title('Box Plot of Odometer Distribution')
plt.show()
```



Notice the distribution is also highly right skewed.

The odometer description shows minimum mileage=0 which is unlikely for second hand cars.

Also the maximum= 10 million which is extremely high.

```
In [49]: # first we check for odometer = 0
df7[(df7['odometer']==0)]
```

Out [49]:

	price	manufacturer	condition	fuel	odometer	transmission	drive	pair
2173	4500	gmc	good	gas	0.0	automatic	rwd	
2269	4250	ford	good	gas	0.0	automatic	rwd	
2355	10500	chevrolet	good	gas	0.0	other	rwd	
2524	7000	ford	good	gas	0.0	automatic	4wd	
5916	9999	jeep	excellent	gas	0.0	automatic	fwd	
...
134351	650	jeep	fair	gas	0.0	automatic	4wd	
137120	20998	mitsubishi	new	gas	0.0	automatic	fwd	
137640	4999	kia	excellent	gas	0.0	automatic	fwd	
139817	6250	chrysler	good	gas	0.0	automatic	rwd	
142597	3500	cadillac	good	gas	0.0	automatic	rwd	

226 rows × 10 columns

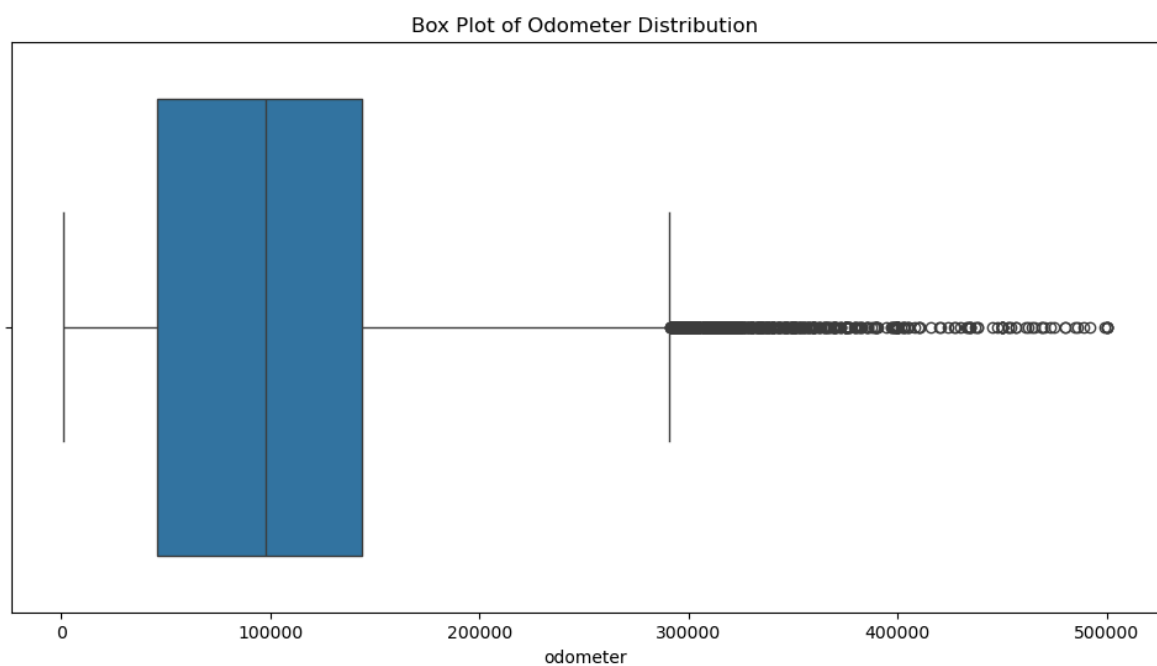

```
In [50]: #even though some of those vehicles with odometer=0 are aged=0, our analy
#hence we will remove all samples with odometer<1000 as anything less tha
#considered a new car
# we will also remove all samples with odometer> 500,000 as they are unma
df8= df7[(df7.odometer>1000) & (df7.odometer<=500000)]

df8.odometer.describe()
```

```
Out [50]: count    141853.000000
mean      100955.105074
std       64332.264191
min       1001.000000
25%       45996.000000
50%       98000.000000
75%       144000.000000
max       500000.000000
Name: odometer, dtype: float64
```

```
In [51]: plt.figure(figsize=(12, 6))
sns.histplot(df8['odometer'], bins=50, kde=True)
plt.title('Histogram of Odometer Distribution')
plt.xlabel('Odometer')
plt.ylabel('Frequency')
plt.show()

plt.figure(figsize=(12, 6))
sns.boxplot(x=df8['odometer'])
plt.title('Box Plot of Odometer Distribution')
plt.show()
```



age of vehicle feature

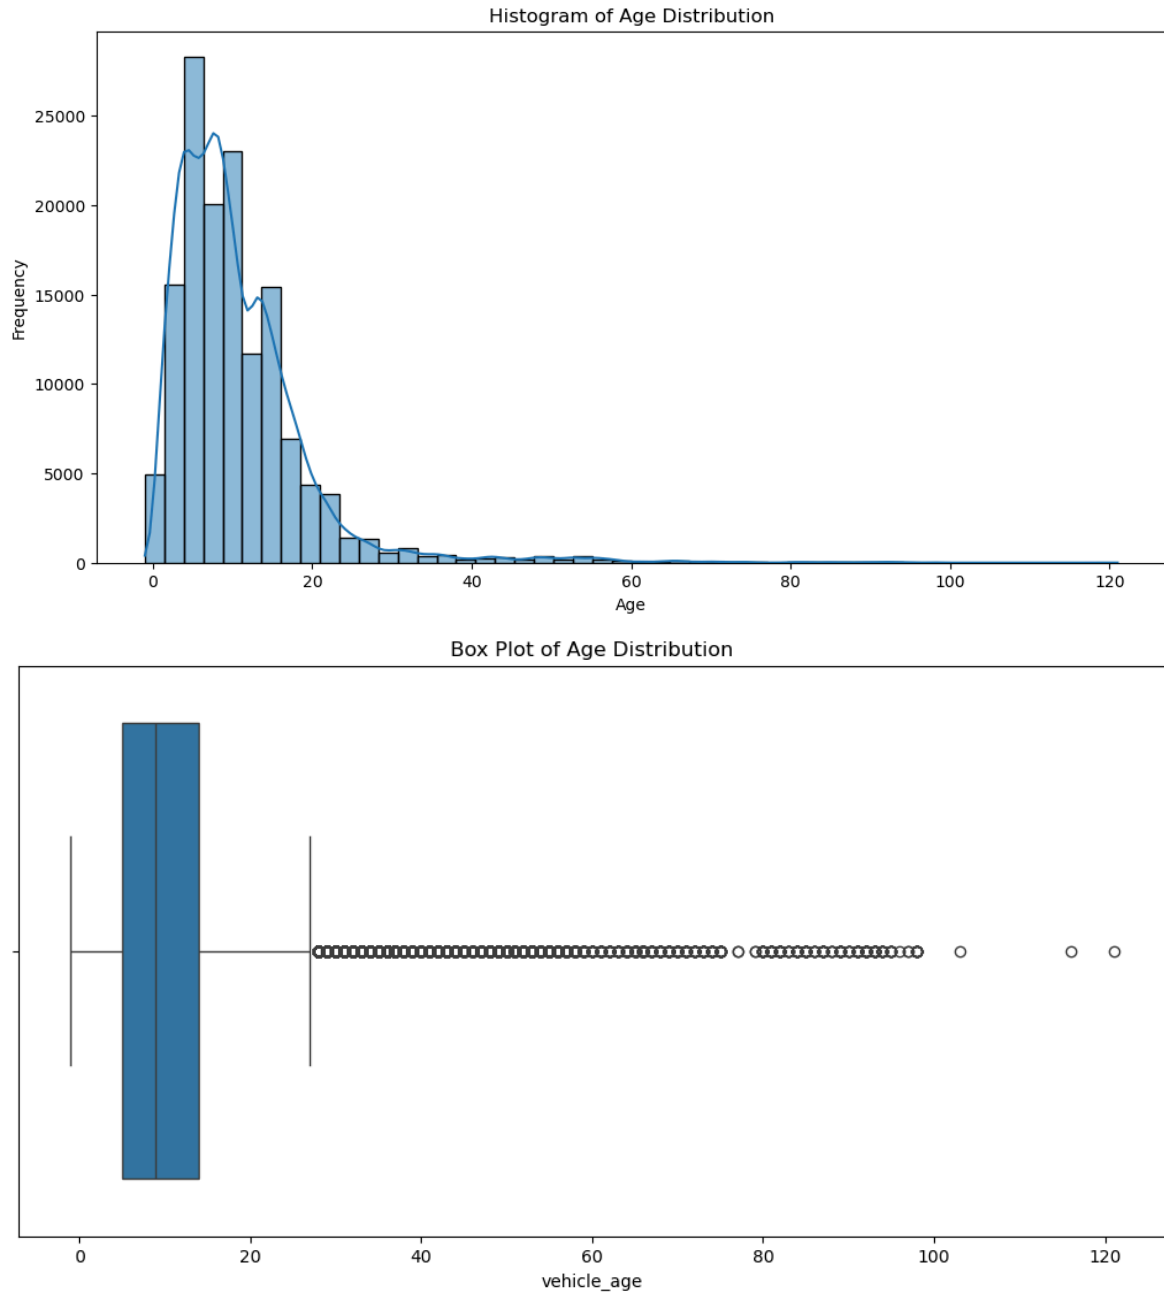
```
In [53]: df8['vehicle_age'].describe(percentiles = [0.25,0.50,0.75,0.85,0.90,1])
```

```
Out[53]: count      141853.000000
mean         10.708184
std           8.918363
min          -1.000000
25%           5.000000
50%           9.000000
75%          14.000000
85%          17.000000
90%          19.000000
100%         121.000000
max          121.000000
Name: vehicle_age, dtype: float64
```

```
In [54]: plt.figure(figsize=(12, 6))
sns.histplot(df8['vehicle_age'], bins=50, kde=True)
plt.title('Histogram of Age Distribution')
```

```
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()

plt.figure(figsize=(12, 6))
sns.boxplot(x=df8['vehicle_age'])
plt.title('Box Plot of Age Distribution')
plt.show()
```



```
In [55]: #remove cars with age<=0 as they can be errors or new cars and hence irre
# we will also split our data into vintage-classic cars (age>=30) and non

#non vintage
df_nonvintage= df8[(df8.vehicle_age>0) & (df8.vehicle_age<30)]

#vintage
df_vintage= df8[(df8.vehicle_age>=30) & (df8.vehicle_age<=100)]

df_nonvintage['vehicle_age'].describe(percentiles = [0.25,0.50,0.75,0.85,
```

```
Out [55]: count      137114.000000
          mean        9.555122
          std         5.826401
          min         1.000000
          25%         5.000000
          50%         8.000000
          75%        13.000000
          85%        16.000000
          90%        18.000000
          100%       29.000000
          max         29.000000
          Name: vehicle_age, dtype: float64
```

Checking for Duplicates

```
In [59]: print('Duplicates in nonvintage data:\n ', df_nonvintage.duplicated().sum()
          '\nDuplicates in vintage data:\n ', df_vintage.duplicated().sum())
```

Duplicates in nonvintage data:
34681

Duplicates in vintage data:
249

```
In [60]: #removing duplicates
df_nonvintage= df_nonvintage.drop_duplicates()
df_vintage= df_vintage.drop_duplicates()

print('Duplicates in nonvintage data:\n ', df_nonvintage.duplicated().sum()
      '\nDuplicates in vintage data:\n ', df_vintage.duplicated().sum())
```

Duplicates in nonvintage data:
0

Duplicates in vintage data:
0

One hot encoding

```
In [67]: #non vintage cars

#to ensure our dummies are not boolean (True/False) set dtype=int

df_nonvintage_encoded= pd.get_dummies(df_nonvintage, columns=['manufacturer',
                                                             'paint_color', 'state'], dtype=int)

df_nonvintage_encoded.head()
```


Out [67]:

	price	odometer	vehicle_age	acura	alfa-romeo	audi	bmw	buick	cadillac	chev
0	55000	167000.0	8	0	0	0	0	0	0	
2	16000	53111.0	3	0	0	0	0	0	0	
4	29000	98000.0	11	0	0	0	0	0	0	
5	23000	94252.0	8	0	0	0	0	0	0	
7	13950	193121.0	14	0	0	0	0	0	0	

5 rows x 120 columns

In [68]: *#to avoide dummy variable trap, we will drop one of the columns from each
#we will choose columns with the least count*

```
#for state= nd (least count), pain_color=custom, fuel= other, transmissio
#condition= salvage (least count), manufacturer= other
df_nonvintage_encoded.drop(['nd','other','custom','other','other','rwd'],'
```

In [70]: *# vintage cars*

```
df_vintage_encoded= pd.get_dummies(df_vintage, columns=['manufacturer', '
                                     'paint_color', 'state'], dtype=in
df_vintage_encoded.drop(['nd','other','custom','other','other','fwd','sal
```

Model

In [72]: *#split into x(input) and y (target/output)*

```
#nonvintage
x_nonvintage=df_nonvintage_encoded.drop('price', axis='columns')
y_nonvintage= df_nonvintage_encoded.price

#vintage
x_vintage=df_vintage_encoded.drop('price', axis='columns')
y_vintage= df_vintage_encoded.price
```

split to train and test

In [76]: **from** sklearn.model_selection **import** train_test_split

```
#since nonvintage cars dataset contains 102,433 samples an 80/20 split wi
x_nonvintage_train, x_nonvintage_test,y_nonvintage_train, y_nonvintage_te

#similarly the vintage cars dataset contains 4,359 samples an 80/20 spli
x_vintage_train, x_vintage_test,y_vintage_train, y_vintage_test= train_te
```

Linear regression

```
In [78]: from sklearn.linear_model import LinearRegression

lr_nonvintage= LinearRegression()
lr_vintage= LinearRegression()

#fit for both vintage and nonvintage
lr_nonvintage.fit(x_nonvintage_train, y_nonvintage_train)
lr_vintage.fit(x_vintage_train, y_vintage_train)
```

```
Out [78]: LinearRegression
LinearRegression()
```

```
In [79]: #evalute the nonvintage model

lr_nonvintage.score(x_nonvintage_test, y_nonvintage_test)
```

```
Out [79]: 0.6904030226735173
```

69% is not bad but still need to improve the model

```
In [81]: #evalute the vintage model

lr_vintage.score(x_vintage_test, y_vintage_test)
```

```
Out [81]: 0.3941757946715474
```

39% is pretty low which is to be expected from vintage cars since we have a smaller sample and it is harder to predict the prices of vintage cars

we will stick to modelling non vintage cars

k-fold Cross Validation

```
In [85]: from sklearn.model_selection import ShuffleSplit
from sklearn.model_selection import cross_val_score

#shufflesplit will randomise the sample to ensure each fold will have equ
#of each of the data samples and is not targeted to one area
cv= ShuffleSplit(n_splits= 5, test_size=0.2, random_state=0)

cross_val_score(LinearRegression(), x_nonvintage, y_nonvintage, cv=cv)
```

```
Out [85]: array([0.68721057, 0.69514816, 0.6933332 , 0.69261292, 0.69100681])
```

In each fold results remain roughly around 69%

Testing different regression models

```
In [88]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Lasso
```

```

from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb

```

```

In [89]: def find_best_model_using_gridsearchcv(X, y):
    algorithms = {
        'linear_regression': {
            'model': LinearRegression(),
            'params': {
                'fit_intercept': [True, False]
            }
        },
        'lasso': {
            'model': Lasso(),
            'params': {
                'alpha': [0.1, 1], # Adjusted alpha values
                'selection': ['random', 'cyclic']
            }
        },
        'decision_tree': {
            'model': DecisionTreeRegressor(),
            'params': {
                'criterion': ['squared_error', 'friedman_mse'],
                'splitter': ['best', 'random']
            }
        }
    },

    scores = []
    cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)

    for algo_name, config in algorithms.items():
        try:
            gs = GridSearchCV(config['model'], config['params'], cv=cv,
                              return_train_score=False, n_jobs=-1)
            gs.fit(X, y)
            scores.append({
                'model': algo_name,
                'best_score_': gs.best_score_,
                'best_params_': gs.best_params_
            })
        except Exception as e:
            print(f"Error in {algo_name}: {e}")

    return pd.DataFrame(scores)

# Call the function with your data
results = find_best_model_using_gridsearchcv(x_nonvintage, y_nonvintage)
results

```

```
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.536e+12, tolerance: 9.964e+08
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.463e+12, tolerance: 9.993e+08
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.542e+12, tolerance: 1.002e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.533e+12, tolerance: 1.001e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.539e+12, tolerance: 1.002e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.861e+10, tolerance: 1.001e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.302e+12, tolerance: 9.964e+08
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.285e+11, tolerance: 1.002e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.017e+12, tolerance: 9.993e+08
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.539e+12, tolerance: 9.964e+08
  model = cd_fast.enet_coordinate_descent(
```

```

/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.535e+12, tolerance: 1.001e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.169e+11, tolerance: 1.002e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.477e+12, tolerance: 1.002e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.544e+12, tolerance: 1.002e+09
  model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear_model/_coordinate_descent.py:697: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.741e+12, tolerance: 1.249e+09
  model = cd_fast.enet_coordinate_descent(

```

Out [89]:

	model	best_score_	best_params_
0	linear_regression	0.691862	{'fit_intercept': True}
1	lasso	0.691769	{'alpha': 0.1, 'selection': 'cyclic'}
2	decision_tree	0.683318	{'criterion': 'friedman_mse', 'splitter': 'ran...

```

In [90]: def find_best_model_using_gridsearchcv(X, y):
          algorithms = {

              'random_forest': {
                  'model': RandomForestRegressor(),
                  'params': {
                      'n_estimators': [50, 100]
                  }
              },

              'xgboost': {
                  'model': xgb.XGBRegressor(),
                  'params': {
                      'n_estimators': [100, 200],
                      'max_depth': [3, 5]
                  }
              }
          }

```

```

    }

    scores = []
    cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)

    for algo_name, config in algorithms.items():
        try:
            gs = GridSearchCV(config['model'], config['params'], cv=cv,
                              return_train_score=False, n_jobs=-1)
            gs.fit(X, y)
            scores.append({
                'model': algo_name,
                'best_score_': gs.best_score_,
                'best_params_': gs.best_params_
            })
        except Exception as e:
            print(f"Error in {algo_name}: {e}")

    return pd.DataFrame(scores)

# Call the function with your data
results = find_best_model_using_gridsearchcv(x_nonvintage, y_nonvintage)
results

```

Out [90]:

	model	best_score_	best_params_
0	random_forest	0.822767	{'n_estimators': 100}
1	xgboost	0.803931	{'max_depth': 5, 'n_estimators': 200}

Random Forest has the best performance with a score of 82%

Random Forest had the best performance

```

In [93]: rf = RandomForestRegressor(n_jobs=-1)
         rf.fit(x_nonvintage_train, y_nonvintage_train)

```

Out [93]:

▼ RandomForestRegressor ⓘ ?

RandomForestRegressor(n_jobs=-1)

```

In [204]: rf.score(x_nonvintage_test, y_nonvintage_test)

```

Out [204]: 0.8257530771207827

```

In [152]: imp=rf.feature_importances_

         col= x_nonvintage_train.columns

         imp_df= pd.DataFrame([imp], columns=col)
         imp_df

```

Out [152...

	odometer	vehicle_age	acura	alfa-romeo	audi	bmw	buick	cadil
0	0.190751	0.418049	0.002148	0.000122	0.001265	0.002702	0.000931	0.00

1 rows × 112 columns

In [202...

```
#we will sum all the levels within each category. since get_dummies arrange in alphabetic order

#manufacturer starts with Acura and ends with Volvo
manufacturer= imp_df.loc[:, 'acura': 'volvo'].sum(axis=1).values[0]

#similarly for the rest of the features
paint_color= imp_df.loc[:, 'black': 'yellow'].sum(axis=1).values[0]
state= imp_df.loc[:, 'ak': 'wy'].sum(axis=1).values[0]
transmission= imp_df.loc[:, 'automatic': 'manual'].sum(axis=1).values[0]
drive= imp_df.loc[:, '4wd': 'fwd'].sum(axis=1).values[0]
fuel= imp_df.loc[:, 'diesel': 'hybrid'].sum(axis=1).values[0]
condition= imp_df.loc[:, 'excellent': 'new'].sum(axis=1).values[0]

imp= pd.DataFrame({ 'manufacturer': [manufacturer], 'paint_color' : [paint_color],
                    'drive': [drive], 'fuel': [fuel], 'condition': [condition] })

#add the first two columns from imp_df to include odometer and vehicle_age
pd.concat([imp, imp_df.iloc[:, :2] ], axis=1)
```

Out [202...

	manufacturer	paint_color	state	transmission	drive	fuel	condition
0	0.087551	0.030071	0.062539	0.014202	0.124431	0.050871	0.021531

We will remove features with less than 5% significance: paint_color, transmission, condition

In [207...

```
#first select columns to be removed
trans= x_nonvintage_train.loc[:, 'automatic': 'manual'].columns.tolist()
color= x_nonvintage_train.loc[:, 'black': 'yellow'].columns.tolist()
cond= x_nonvintage_train.loc[:, 'excellent': 'new'].columns.tolist()

#combine those columns
col_remove= trans+color+cond

#modify both train and test sets

x_nonvintage_train_reduced =x_nonvintage_train.drop(columns= col_remove)
x_nonvintage_test_reduced= x_nonvintage_test.drop(columns= col_remove)
y_nonvintage_train_reduced =y_nonvintage_train.drop(columns= col_remove)
y_nonvintage_test_reduced= y_nonvintage_test.drop(columns= col_remove)
```

Implement Random Forest on the reduced data

In [253...

```
rf_reduced= RandomForestRegressor(n_jobs=-1)
rf_reduced.fit(x_nonvintage_train_reduced, y_nonvintage_train_reduced)
```

Out [253...

▼ RandomForestRegressor ⓘ ↻

RandomForestRegressor(n_jobs=-1)

In [254...

```
rf_reduced.score(x_nonvintage_test_reduced, y_nonvintage_test_reduced)
```

Out [254...

```
0.8046099029041316
```

Model performance did not improve. try with cross validation

In [216...

```
x_nonvintage_reduced = x_nonvintage.drop(columns= col_remove)
y_nonvintage_reduced = y_nonvintage.drop(columns= col_remove)
```

In [229...

```
cv= ShuffleSplit(n_splits= 5, test_size=0.2, random_state=0)

cross_val_score(RandomForestRegressor(n_jobs=-1), x_nonvintage_reduced, y
```

Out [229...

```
array([0.80388894, 0.80362325, 0.80208947, 0.79945176, 0.80429128])
```

In [230...

```
# try with the second best model: xgboost

cv= ShuffleSplit(n_splits= 5, test_size=0.2, random_state=0)

cross_val_score(xgb.XGBRegressor(n_jobs=-1), x_nonvintage_reduced, y_nonv
```

Out [230...

```
array([0.78439391, 0.78865314, 0.78627592, 0.78047609, 0.78332841])
```

Removing features did not improve our model. We will keep all features

In [233...

```
rf= RandomForestRegressor(n_jobs=-1)
rf.fit(x_nonvintage_train, y_nonvintage_train)
rf.score(x_nonvintage_test, y_nonvintage_test)
```

Out [233...

```
0.8258058553370891
```

In [305...

```
def predict_price(manufacturer, condition, fuel, odometer, transmission,
# Initialize an array of the size of all the columns in x_nonvintage
x = np.zeros(len(x_nonvintage.columns))

#set first two indices to our numeric features
x[0] = odometer
x[1] = vehicle_age

# list of categorical features
features = [state, manufacturer, condition, fuel, transmission, paint

# for all categories in a categorical feature, loop over the categorie
for category in features:

    # Check if the category is in our data x_nonvintage
    if category in x_nonvintage.columns:
```



```

        # Get the index for the one hot encoded category in that data
        idx = np.where(x_nonvintage.columns == category)[0][0]

        #set that category index=1 (categories for that categorical f
        x[idx] = 1

    # prediction
    return rf.predict([x])[0]

```

In [307... predict_price('bmw', 'like new', 'hybrid', 120000, 'automatic', 'fwd', 'gr

```

/opt/anaconda3/lib/python3.12/site-packages/sklearn/base.py:493: UserWarni
ng: X does not have valid feature names, but RandomForestRegressor was fit
ted with feature names
    warnings.warn(

```

Out[307... 9723.135

In [309... predict_price('bmw', 'like new', 'hybrid', 100000, 'automatic', 'fwd', 'gr

```

/opt/anaconda3/lib/python3.12/site-packages/sklearn/base.py:493: UserWarni
ng: X does not have valid feature names, but RandomForestRegressor was fit
ted with feature names
    warnings.warn(

```

Out[309... 15187.78

In [311... predict_price('bmw', 'good', 'hybrid', 100000, 'hybrid', 'fwd', 'grey', 'ny

```

/opt/anaconda3/lib/python3.12/site-packages/sklearn/base.py:493: UserWarni
ng: X does not have valid feature names, but RandomForestRegressor was fit
ted with feature names
    warnings.warn(

```

Out[311... 16578.1

```

In [320... import pickle
with open('/Users/marwa/Desktop/2ndHCP/model/rf_model', 'wb') as f:
    pickle.dump(rf, f)

```

```

In [321... import json
columns = {
    'data_columns' : [col.lower() for col in x_nonvintage.columns]
}

with open('/Users/marwa/Desktop/2ndHCP/model/columns', 'w') as f:
    f.write(json.dumps(columns))

```